



Concurrency



Agenda

- WHAT THE CONCURRENCY IS?
- DATA SYNCHRONIZATION
- DATA RACE




Concurrency vs parallelism

Single Threaded



Parallel



Concurrent





What is a PROCESS?

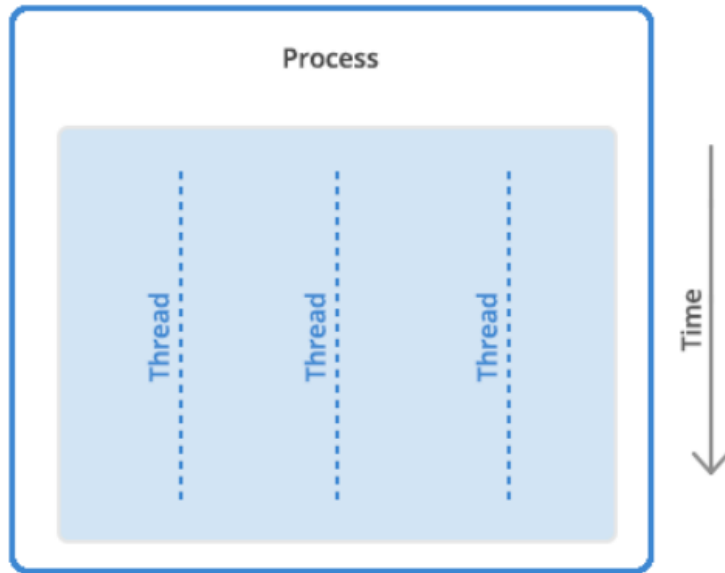
A process is the execution of a program that allows you to perform the appropriate actions specified in a program. It can be defined as an execution unit where a program runs. The OS helps you to create, schedule, and terminates the processes which is used by CPU. The other processes created by the main process are called child process.





What is a THREAD?

Thread is an execution unit that is part of a process. A process can have multiple threads, all executing at the same time.



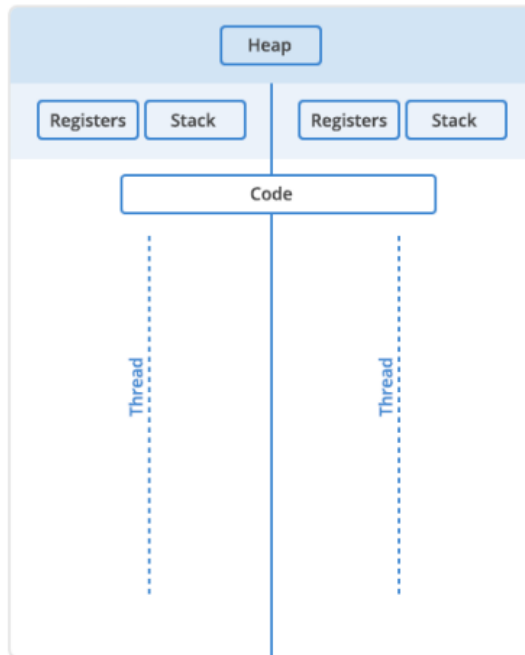


Single/Multi Thread Process

Single Thread



Multi Threaded






Process vs Thread

PROCESS	THREAD
Processes are heavyweight operations	Threads are lighter weight operations
Each process has its own memory space	Threads use the memory of the process they belong to
Inter-process communication is slow as processes have different memory addresses	Inter-thread communication can be faster than inter-process communication because threads of the same process share memory with the process they belong to
Context switching between processes is more expensive	Context switching between threads of the same process is less expensive
Processes don't share memory with other processes	Threads share memory with other threads of the same process



Parallel Execution


Parallel execution is when a computer has more than one CPU or CPU core, and makes progress on more than one task simultaneously. However, parallel execution is not referring to the same phenomenon as parallelism.





Parallel Execution


Concurrency means that an application is making progress on more than one task - at the same time or at least seemingly at the same time (concurrently). If the computer only has one CPU the application may not make progress on more than one task at exactly the same time, but more than one task is in progress at a time inside the application. To make progress on more than one task concurrently the CPU switches between the different tasks during execution.





Parallel Concurrent Execution

It is possible to have parallel concurrent execution, where threads are distributed among multiple CPUs. Thus, the threads executed on the same CPU are executed concurrently, whereas threads executed on different CPUs are executed in parallel. The diagram below illustrates parallel concurrent execution.





Goroutine



Goroutines

- What is a goroutine? It's an independently executing function, launched by a go statement.
- It has its own call stack, which grows and shrinks as required.
- It's very cheap. It's practical to have thousands, even hundreds of thousands of goroutines.
- It's not a thread.
- There might be only one thread in a program with thousands of goroutines.
- Instead, goroutines are multiplexed dynamically onto threads.



GOROUTINE

A goroutine is a lightweight thread managed by the Go runtime.

```
go f(x, y, z)
```

starts a new goroutine

```
f(x, y, z)
```

The evaluation of `f`, `x`, `y`, and `z` happens in the current goroutine and the execution of `f` happens in the new goroutine.

Goroutines run in the same address space, so access to shared memory must be synchronized.

GOROUTINE EXAMPLE



code

```
package main

import (
    "fmt"
)

func main() {
    go fmt.Println("world")
    fmt.Println("hello")
}
```

output

hello

<https://goplay.space/#EUdGJ0HsYcf>



Channels



CHANNELS

Channels are a typed channel through which you can send and receive data with the channel operator, `<-`.

Like maps, channels must be created before use:

```
ch := make(chan int)
```

```
ch <- v    // Send v to channel ch.
```

```
v := <-ch  // Receive from ch, and  
           // assign value to v.
```

(The data flows in the direction of the arrow.)

By default, sends and receives block until the other side is ready. This allows goroutines to synchronize without explicit locks or condition variables.



Channels (Initialization)

code

```
func main() {  
    ch := make(chan string, 1)  
  
    go func() {  
        ch <- "Hello World"  
    }()  
  
    fmt.Println(<-ch)  
}
```

output

Hello World

Channel initialization:

make(chan type, n)

type - type of the value

n - length of the inner buffer

<https://goplay.space/#CfKMQLhGFIZ>



Channels (Initialization)

code

```
func main() {  
    var ch chan string  
  
    go func() {  
        ch <- "Hello World"  
    }()  
  
    fmt.Println(<-ch)  
}
```

output

```
fatal error: all goroutines are asleep - deadlock!  
goroutine 1 [chan receive (nil chan)]: main.main()  
/tmp/sandbox1008553871/prog.go:14 +0x4e goroutine 6  
[chan send (nil chan)]: main.main.func1()  
/tmp/sandbox1008553871/prog.go:11 +0x25 created by  
main.main /tmp/sandbox1008553871/prog.go:10 +0x3c
```

Read and Write from channel that isn't initialized with, bring lock to the go routine.

<https://goplay.space/#Gz4z9gcKZw7>



CHANNELS EXAMPLE

code

```
func main() {  
    c := make(chan int)  
    c <- 2  
}
```

output

fatal error: all goroutines are asleep - deadlock!

goroutine 1 [chan send]:
main.main()
C:/ttt/main.go:5 +0x31



Channels (Read & Write)

code

```
func main() {  
    ch := make(chan string)  
  
    go func() {  
        ch <- "Hello World"  
    }()  
  
    fmt.Println(<-ch)  
}
```

output

Hello World

Communication between Go routines
Insert value into channel
ch <- "Hello World"

Read from channel
x := <- ch

<https://goplay.space/#Dq31pd30U62>



Range and Close

A sender can close a channel to indicate that no more values will be sent. Receivers can test whether a channel has been closed by assigning a second parameter to the receive expression:
After

```
v, ok := <-ch
```

ok is false if there are no more values to receive and the channel is closed.

The loop

```
for i := range c
```

receives values from the channel repeatedly until it is closed.

Note: Only the sender should close a channel, never the receiver. Sending on a closed channel will cause a panic.

Another note: Channels aren't like files; you don't usually need to close them. Closing is only necessary when the receiver must be told there are no more values coming, such as to terminate a range loop.



Channels (Finalization)

code

```
func main() {  
    ch := make(chan string)  
  
    go func() {  
        ch <- "Hello World"  
        close(ch)  
    }()  
  
    for m := range ch {  
        fmt.Println(m)  
    }  
}
```

output

Hello World

Channel finalization:

close(ch)

ch – name of the variable

<https://goplay.space/#JsAiRyH6TID>



Channels(Finalization)

code

```
func main() {  
    ch := make(chan string)  
  
    close(ch)  
  
    ch <- "Hello World"  
}
```

output

```
panic: send on closed channel  
goroutine 1 [running]: main.main()  
/tmp/sandbox2454463867/prog.go:8  
+0x45
```

Write to closed channel cause panic.

<https://goplay.space/#CwM7VmJ2YLJ>



Channels (Finalization)

code

```
func main() {  
    ch := make(chan string)  
  
    go func() {  
        ch <- "Hello World"  
        close(ch)  
    }()  
  
    for i := 0; i < 3; i++ {  
        v, ok := <-ch  
        fmt.Printf("Value - %q, Ok - %t\n", v, ok)  
    }  
}
```

output

```
Value - "Hello World", Ok - true  
Value - "", Ok - false  
Value - "", Ok - false
```

Read from closed channel always return default value for the type and flag with value "false".

<https://goplay.space/#DmjdkGkmNZc>



Range and Close

code

```
func fibonacci(n int, c chan int) {  
    x, y := 0, 1  
    for i := 0; i < n; i++ {  
        c <- x  
        x, y = y, x+y  
    }  
    close(c)  
}  
func main() {  
    c := make(chan int, 10)  
    go fibonacci(cap(c), c)  
    for i := range c {  
        fmt.Println(i)  
    }  
}
```

output

```
0  
1  
1  
2  
3  
5  
8  
13  
21  
34
```

range over channel

Loop that range over channel runs until channel closing.

https://goplay.space/#INbU_AM_z_O



CHANNELS EXAMPLE

code

```
func sum(s []int, c chan int) {  
    sum := 0  
    for _, v := range s {  
        sum += v  
    }  
    c <- sum // send sum to c  
}
```

```
func main() {  
    s := []int{7, 2, 8, -9, 4, 0}  
    c := make(chan int)  
    go sum(s[:len(s)/2], c)  
    go sum(s[len(s)/2:], c)  
    x, y := <-c, <-c // receive from c  
    fmt.Println(x, y, x+y)  
}
```

output

-5 17 12

<https://goplay.space/#aCVqbyM4aDq>



BUFFERED CHANNELS

Channels can be buffered. Provide the buffer length as the second argument to make to initialize a buffered channel:

```
ch := make(chan int, 100)
```

Sends to a buffered channel block only when the buffer is full. Receives block when the buffer is empty.

code

```
func main() {  
    ch := make(chan int, 2)  
    ch <- 1  
    ch <- 2  
    fmt.Println(<-ch)  
    fmt.Println(<-ch)  
}
```

output

```
1  
2
```



BUFFERED CHANNELS

code

```
func main() {  
    ch := make(chan string, 10)  
    fmt.Println(len(ch), cap(ch))  
    ch <- "Hello World"  
    fmt.Println(len(ch), cap(ch))  
    ch <- "Hello World"  
    fmt.Println(len(ch), cap(ch))  
}
```

output

```
0 10  
1 10  
2 10
```



Select



Select

code

```
func fibonacci(c, quit chan int) {  
    x, y := 0, 1  
    for {  
        select {  
        case c <- x:  
            x, y = y, x+y  
        case <-quit:  
            fmt.Println("quit")  
            return  
        }  
    }  
}  
  
func main() {  
    c := make(chan int)  
    quit := make(chan int)  
    go func() {  
        for i := 0; i < 10; i++ {  
            fmt.Println(<-c)  
        }  
        quit <- 0  
    }()  
    fibonacci(c, quit)  
}
```

output

```
0  
1  
1  
2  
3  
5  
8  
13  
21  
34  
quit
```

<https://goplay.space/#eJZ6RZrnwcA>



Default section

code

```
func main() {  
    tick := time.Tick(100 * time.Millisecond)  
    boom := time.After(500 * time.Millisecond)  
    for {  
        select {  
        case <-tick:  
            fmt.Println("tick.")  
        case <-boom:  
            fmt.Println("BOOM!")  
        return  
        default:  
            fmt.Println(".")  
            time.Sleep(50 * time.Millisecond)  
        }  
    }  
}
```

output

```
.  
.  
tick.  
.  
.  
tick.  
.  
.  
tick.  
.  
BOOM!
```

<https://goplay.space/#7gu7TzopGIM>



Sync



Mutex

A Mutex, or a mutual exclusion is a mechanism that allows us to prevent concurrent processes from entering a critical section of data whilst it's already being executed by a given process.

Race condition

```
go func() {  
    i++  
}()  
go func() {  
    i++  
}()
```



```
tmp = a+1  
a = tmp  
if a == 1 {  
    criticalSection()  
}
```

```
tmp = a+1  
a = tmp  
if a == 1 {  
    criticalSection()  
}
```

```
var mx sync.Mutex  
// ..  
mx.Lock()  
tmp = a + 1  
a = tmp  
if a == 1 {  
    criticalSection()  
}  
mx.Unlock()
```



Sync: WaitGroup

code

```
func worker(id int, wg *sync.WaitGroup) {  
    defer wg.Done()  
    fmt.Printf("Worker %d starting\n", id)  
    time.Sleep(time.Second)  
    fmt.Printf("Worker %d done\n", id)  
}  
  
func main() {  
    var wg sync.WaitGroup  
    for i := 1; i <= 5; i++ {  
        wg.Add(1)  
        go worker(i, &wg)  
    }  
    wg.Wait()  
}
```

output

```
Worker 2 starting  
Worker 3 starting  
Worker 4 starting  
Worker 5 starting  
Worker 1 starting  
Worker 1 done  
Worker 4 done  
Worker 3 done  
Worker 2 done  
Worker 5 done
```

https://goplay.space/#_5jBUZmYcFz



Data Race



What is it?

Data races are among the most common and hardest to debug types of bugs in concurrent systems. A data race occurs when two goroutines access the same variable concurrently and at least one of the accesses is a write.

code

```
func race() {  
    wait := make(chan struct{})  
    n := 0  
    go func() {  
        n++ // read, increment, write  
        close(wait)  
    }()  
    n++ // conflicting access  
    <-wait  
    fmt.Println(n) // Output: <unspecified>  
}
```



How to find?

To help diagnose such bugs, Go includes a built-in data race detector. To use it, add the `-race` flag to the `go` command:

```
$ go test -race mypkg      // to test the package
$ go run -race mysrc.go    // to run the source file
$ go build -race mycmd     // to build the command
$ go install -race mypkg   // to install the package
```



Typical Data Races: Race on loop counter

code

```
func main() {  
    var wg sync.WaitGroup  
    wg.Add(5)  
    for i := 0; i < 5; i++ {  
        go func() {  
            fmt.Println(i)  
            wg.Done()  
        }()  
    }  
    wg.Wait()  
}
```

output

```
3  
5  
5  
5  
5
```

fix

```
func main() {  
    var wg sync.WaitGroup  
    wg.Add(5)  
    for i := 0; i < 5; i++ {  
        go func(j int) {  
            fmt.Println(j)  
            wg.Done()  
        }(i)  
    }  
    wg.Wait()  
}
```

https://goplay.space/#_qA-EoVRpJw



Typical Data Races: Accidentally shared variable

code

```
// ParallelWrite writes data to file1 and file2, returns the errors.
func ParallelWrite(data []byte) chan error {
    res := make(chan error, 2)
    f1, err := os.Create("file1")
    if err != nil {
        res <- err
    } else {
        go func() {
            // This err is shared with the main goroutine,
            // so the write races with the write below.
            _, err = f1.Write(data)
            res <- err
            f1.Close()
        }()
    }
    f2, err := os.Create("file2") // The second conflicting write to err.
    if err != nil {
        res <- err
    } else {
        go func() {
            _, err = f2.Write(data)
            res <- err
            f2.Close()
        }()
    }
}
```

fix

```
...
_, err := f1.Write(data)
...
_, err := f2.Write(data)
...
```



Typical Data Races: Unprotected global variable

code

```
var service map[string]net.Addr

func RegisterService(name string, addr net.Addr) {
    service[name] = addr
}

func LookupService(name string) net.Addr {
    return service[name]
}
```

fix

```
var (
    service map[string]net.Addr
    serviceMu sync.Mutex
)

func RegisterService(name string, addr net.Addr) {
    serviceMu.Lock()
    defer serviceMu.Unlock()
    service[name] = addr
}

func LookupService(name string) net.Addr {
    serviceMu.Lock()
    defer serviceMu.Unlock()
    return service[name]
}
```




Atomic counters

code

```
func main() {  
    var ops uint64  
    var wg sync.WaitGroup  
    for i := 0; i < 50; i++ {  
        wg.Add(1)  
        go func() {  
            for c := 0; c < 1000; c++ {  
                ops++  
            }  
            wg.Done()  
        }()  
    }  
    wg.Wait()  
    fmt.Println("ops:", ops)  
}
```

output

ops: 43518

<https://goplay.space/#Vy5J5SIOBMw>



Atomic counters fix

code

```
func main() {  
    var ops uint64  
    var wg sync.WaitGroup  
    for i := 0; i < 50; i++ {  
        wg.Add(1)  
        go func() {  
            for c := 0; c < 1000; c++ {  
                atomic.AddUint64(&ops, 1)  
            }  
            wg.Done()  
        }()  
    }  
    wg.Wait()  
    fmt.Println("ops:", ops)  
}
```

output

ops: 50000

<https://goplay.space/#-v842yznjla>



Mutex counter

code

```
type SafeCounter struct {  
    mu sync.Mutex  
    v  map[string]int  
}  
  
func (c *SafeCounter) Inc(key string) {  
    c.mu.Lock()  
    c.v[key]++  
    c.mu.Unlock()  
}  
  
func (c *SafeCounter) Value(key string) int {  
    c.mu.Lock()  
    defer c.mu.Unlock()  
    return c.v[key]  
}  
  
func main() {  
    c := SafeCounter{v: make(map[string]int)}  
    for i := 0; i < 1000; i++ {  
        go c.Inc("somekey")  
    }  
    time.Sleep(time.Second)  
    fmt.Println(c.Value("somekey"))  
}
```

output

1000

https://goplay.space/#UUh_4XFIJtT



Context



Cancel go routine with context

code

```
func main() {  
    work := make(chan func(), 10)  
    ctx, cancel := context.WithCancel(context.Background())  
    var wg sync.WaitGroup  
    wg.Add(1)  
    go func() {  
        for {  
            select {  
            case <-ctx.Done():  
                fmt.Println("Worker shutdown")  
                wg.Done()  
                return  
            case w := <-work:  
                w()  
            }  
        }  
    }()  
    for i := 0; i < 10; i++ {  
        work <- func() {  
            fmt.Println("Working...")  
            time.Sleep(1 * time.Second)  
        }  
    }  
    <-time.After(2 * time.Second)  
    cancel()  
    wg.Wait()  
}
```

output

```
T+00.000000Working...  
T+01.000000Working...  
T+02.000000Working...  
T+03.000000Working...  
T+04.000000Worker shutdown
```

<https://goplay.space/#TuU9qqWyhpX>



Handle timeout & Deadlines

code

output

T+00.002000context deadline exceeded

```
const shortDuration = 2 * time.Millisecond
```

```
func main() {  
    ctx, cancel := context.WithTimeout(context.Background(), shortDuration)  
    defer cancel()
```

```
    // d := time.Now().Add(shortDuration)  
    // ctx, cancel := context.WithDeadline(context.Background(), d)
```

```
    select {  
    case <-time.After(1 * time.Second):  
        fmt.Println("overslept")  
    case <-ctx.Done():  
        fmt.Println(ctx.Err()) // prints "context deadline exceeded"  
    }  
}
```

<https://goplay.space/#-5E6JKWdSyx>



ErrGroup



Semaphore (Limit parallelism for task)

code

```
func main() {  
    const paralel = 5  
    var wg sync.WaitGroup  
    sem := make(chan struct{}, paralel)  
    for i := 0; i < 100; i++ {  
        wg.Add(1)  
        sem <- struct{}{}  
        go func() {  
            work()  
            <-sem  
            wg.Done()  
        }()  
    }  
    wg.Wait()  
}  
  
func work() {  
    time.Sleep(1 * time.Second)  
    fmt.Println("Work is done")  
}
```

output

```
T+01.000000 Work is done  
Work is done  
Work is done  
Work is done  
Work is done  
T+02.000000 Work is done  
Work is done  
Work is done
```

<https://goplay.space/#Q8txT0BL5kJ>



Semaphore (Limit parallelism for task) errgroup

code

```
func main() {  
    const parallel = 5  
    errG, _ := errgroup.WithContext(context.Background())  
    errG.SetLimit(parallel)  
    for i := 0; i < 100; i++ {  
        errG.Go(func() error {  
            return work()  
        })  
    }  
    err := errG.Wait()  
}  
  
func work() error {  
    time.Sleep(1 * time.Second)  
    fmt.Println("Work is done")  
    return nil  
}
```

output

```
T+01.000000 Work is done  
Work is done  
Work is done  
Work is done  
Work is done  
T+02.000000 Work is done  
Work is done  
Work is done
```

<https://goplay.space/#5Ci08VNbqM->



Questions



Additional materials

Links

- https://golang.org/doc/effective_go.html#concurrency
- https://www.youtube.com/watch?v=cN_DpYBzKso
- <https://github.com/golang/go/wiki/LearnConcurrency>
- <https://habr.com/ru/company/avito/blog/466495/>



Task

Homework

Concurrent batch

- Implement a function that will load users from the database

```
getBatch(n int64, pool int64) res []user
```

- The function takes two arguments - the number of users and the number of goroutines in which users will concurrently load

```
for i := 0; i < number; i++ {
```

```
    go func() {
```

```
        user := getOne...
```

- The function returns an array of received users
- It is planned to review the solution for the interview

Tips & tricks

- Don't forget the data race
- Due to autocode VM limitations, run tests locally. It is possible that autocode tests will pass with an invalid solution
- To limit concurrently running goroutines you can use:
 - ☐ Semaphore pattern
 - ☐ Worker pool
 - ☐ errgroup
 - ☐ etc



Thanks